

# The experiments are complete

## Trustworthy Evaluation Frameworks for AI-Assisted Painting Restoration

This brief replaces the old expansion plan as the current meeting handout. It reports completed Controlled-300 evidence, not projected workloads. No additional numbered notebook is required for the accepted closeout.

| Completed evidence       | Scope                                                           |
|--------------------------|-----------------------------------------------------------------|
| Collection               | 300 paintings / five balanced visual categories                 |
| Experimental registry    | 3,425 cases; 2,620 restoration-eligible                         |
| Matched comparison       | 10,480 primary candidates across four full methods              |
| Broader report catalogue | 13,879 retained candidates, separate from all execution records |
| Repeated-seed evidence   | 1,025 supported four-seed groups                                |
| SDXL feasibility         | 24 completed candidates from 35 scheduled cases                 |
| N36 delivery             | 134 files, including 123 package files                          |
| N36 validation           | 859 passed / 16 warning nonpasses / zero blockers               |

**Main result:** LaMa leads 10 of 11 saved quality anchors; Telea leads crop SSIM. The contribution is the evidence framework that preserves this disagreement, not a universal model ranking.

[Open the completed dashboard](#)

Evidence baseline: committed closeout 453b4bfa. Scientific source: N36 supervisor summary, key findings, run manifest and reproducibility appendix. This post-closeout handout does not alter those frozen artifacts.

## What remains is interpretation and delivery

### Questions for the supervisor

1. Confirm that the final RQ answers and their limits are appropriate for the thesis.
2. Select the principal figures and representative case narratives.
3. Agree the emphasis of SDXL feasibility and the separate D01/D02 studies.
4. Review the limitations and distinguish completed work from possible future research.

### Work that is not another experiment

Thesis writing and revision; defence preparation; final editorial review; optional repository housekeeping; and a publication/archive decision. Zenodo is planned, not published. None is a reason to silently modify the completed scientific record.

### Disclosed limitations remain part of the result

The N36 warning status preserves four dependency warnings, ten unmeasured deployment qualifications and two provenance exceptions (N29 historical manifest; N35 notebook-source mismatch). Successful availability runs and owner acceptance do not replace unmeasured performance or independent reproduction evidence.

### No conservation approval claim

Synthetic damage differs from physical conservation. Metrics do not establish authenticity, repeated-seed disagreement is not calibrated confidence, and computational flags are not expert judgements. Rendered lightness in D02 is not demographic identity.

#### **Recommended closing message:**

"The implementation and experimentation are complete. We can now focus on how clearly the thesis explains the evidence, the disagreements and the boundaries of its conclusions."

[Final supervisor synthesis](#) / [Reproducibility appendix](#) / [Limitations](#)